

Multi-Agent Mamba-RL
Pipeline

MediaTek\_2026\_project

2026 8

Abstract

MediaTek\_2026\_project Multi-Agent Mamba-RL
Mamba-2.8B
agent agent coordinator agent agent
LoRA agent Mamba selective state-space
recurrence coordinator specialist
coordinator agent REINFORCE MovieLens Amazon
Reviews 2023 Yelp 2019/2023
Mamba loss gradient

Contents

# 1

pipeline src/train.py src/train\_rl.py

1. agent agent coordinator
2. agent LoRA LoRA specialist
3. Mamba 2.8B Mamba RL update selective recurrence
4. specialist joint RL
5. train/validation/test catalog
6. agent grounded prompt

# 2

pipeline ?? launcher import

Table 1: Multi-Agent Mamba-RL

| run_mamba_rl.sh        | GPU cache                                   |
|------------------------|---------------------------------------------|
| src/data_mamba_rl.py   | MovieLens Amazon Yelp synthetic             |
| src/model_mamba_rl.py  | LoRA selective-state specialist coordinator |
| src/train_mamba_rl.py  | REINFORCE                                   |
| tests/test_mamba_rl.py | agent stage shape                           |
| outputs_mamba_rl/      | dataset/run ID checkpoint metrics           |

# 3 Pipeline

?? full-catalog ranking




Figure 1: Multi-Agent Mamba-RL

### 3.1

Input:  $i$  (title, genre, category, feature, description, business metadata),  $t_i$  (Mamba)

$$e_i^{\text{text}} = \text{MeanPool}(\text{Mamba}_{\text{frozen}}(t_i)) \in \mathbb{R}^{d_m}. \quad (1)$$

Cache:  $E_{\text{text}} \in \mathbb{R}^{N \times d_m}$ . Run:  $E_{\text{text}}$  (2.8B)

$$e_i = \text{norm}(W_p e_i^{\text{text}}), \quad W_p \in \mathbb{R}^{d \times d_m}, \quad (2)$$

Agent-specific adaptation:  $d = 128$ ,  $W_p$  (agent-specific LoRA)

## 4

### 4.1

loader

$$(u, i, \tau, t_i), \quad (3)$$

Input:  $u$  (ID),  $i$  (ID),  $\tau$  (timestamp),  $t_i$  (ID). MovieLens, Yelp, rating/stars  $\geq 4$ , implicit feedback.

### 4.2

Table 2:

| dataset                                    | features                   | metadata                                                        |
|--------------------------------------------|----------------------------|-----------------------------------------------------------------|
| movielens-100k/1m/20m/25m/32m/latest-small | rating, timestamp          | title, genres, GroupLens                                        |
| amazon-<category>                          | Amazon Reviews 2023 review | metadata, title, subtitle, main category, features, description |

| □□□□□                     | □□□□                                   | □□□□                                                    |
|---------------------------|----------------------------------------|---------------------------------------------------------|
| amazon:<exact-<br>config> | □□ Hugging Face raw re-<br>view config | □□□ raw meta config □ parent_asin<br>join□              |
| yelp19, yelp23            | Yelp review JSON □□□□<br>□□            | business name□categories□city□<br>state □□□ attributes□ |
| synthetic                 | □□□□□□□□□□                             | □□□□□□□□□□ smoke test□                                  |

Yelp □□ snapshot □□□□□□□□□□ launcher □□□□□□□□□□ DATA\_PATH □□□□  
 □□□CSV□JSONL □ Parquet□□□ Yelp JSON □□□ yelp\_academic\_dataset\_review.json  
 □ yelp\_academic\_dataset\_business.json□

### 4.3 Chronological split

□□□□  $u$  □□□□□□□□□□

$$\mathcal{H}_u = (i_1, i_2, \dots, i_{T_u}). \quad (4)$$

□□ leave-two-out□

$$\mathcal{H}_u^{\text{train}} = (i_1, \dots, i_{T_u-2}), \quad (5)$$

$$i_u^{\text{valid}} = i_{T_u-1}, \quad (6)$$

$$i_u^{\text{test}} = i_{T_u}. \quad (7)$$

Validation □□□□ training history□test □□□□ training history □ validation item□□□□□□  
 □□□□□□□ state□

### 4.4 RL transition

training history □□□□ prefix-to-next-item transition□

$$((i_1, \dots, i_k), i_{k+1}), \quad k = 1, \dots, T_u - 3. \quad (8)$$

□ transition □□□□ MAX\_TRANSITIONS□□□ reservoir sampling □□□□□□□□□□ MovieLens-  
 20M □□□ Amazon/Yelp □□□□□□□□□□□□□ transition list□

## 5 Multi-Agent □□□□

### 5.1 Agent □□

□□□□□□□ agent □□□

#### Long-term Agent

□□□□ max-history □□□□□□□□□□□□ time scale□0.08□□□□□□□□□□  
 □□□□□

#### Short-term Agent

□□□□□ short-window □□□□□□ 10□□□□□□ time scale□0.5□□□□□□□□□  
 session intent□

## Coordinator Agent

state mixing weights specialist scores  
 coordinator score

agent LoRA agent GPU  
 Mamba-2.8B agent specialization

## 5.2 LoRA

$\mathbf{x} \in \mathbb{R}^{d_{in}}$  agent-specific LoRA update

$$\text{LoRA}_j(\mathbf{x}) = \frac{\alpha}{r} B_j A_j \text{Dropout}(\mathbf{x}), \quad (9)$$

$A_j \in \mathbb{R}^{r \times d_{in}}$   $B_j \in \mathbb{R}^{d_{out} \times r}$   $r = 8$   $\alpha = 16$  dropout = 0.05  $A_j$  Kaiming  
 initialization  $B_j$  adapter base path

agent candidate delta gate output LoRA coordinator  
 state mix score LoRA agent adapter

## 5.3 Mamba selective state update

$t$   $\mathbf{x}_t$  specialist  $j \in \{L, S\}$

$$\delta_t^{(j)} = \text{softplus}(\mathbf{b}_\delta^{(j)} + \text{LoRA}_\delta^{(j)}(\mathbf{x}_t)), \quad (10)$$

$$\mathbf{a}_t^{(j)} = \exp(-\delta_t^{(j)}), \quad (11)$$

$$\tilde{\mathbf{h}}_t^{(j)} = \tanh(\mathbf{x}_t + \text{LoRA}_{\text{candidate}}^{(j)}(\mathbf{x}_t)), \quad (12)$$

$$\mathbf{h}_t^{(j)} = \mathbf{a}_t^{(j)} \odot \mathbf{h}_{t-1}^{(j)} + (1 - \mathbf{a}_t^{(j)}) \odot \tilde{\mathbf{h}}_t^{(j)}, \quad (13)$$

$$\mathbf{g}_t^{(j)} = \text{sigmoid}(\text{LoRA}_{\text{gate}}^{(j)}(\mathbf{x}_t)), \quad (14)$$

$$\mathbf{o}_t^{(j)} = \mathbf{g}_t^{(j)} \odot \mathbf{h}_t^{(j)} + (1 - \mathbf{g}_t^{(j)}) \odot \mathbf{x}_t. \quad (15)$$

output LoRA  $\ell_2$  normalization recurrence  $T$   $\mathcal{O}(Td)$   
 Transformer  $T \times T$  attention matrix

$\delta_t$  input-sensitive forgetting rate state  
 long/short agents time scale observation window specialization  
 inductive bias loss

## 5.4 Specialist policy scores

$\mathbf{e}_i$

$$z_i^{(j)} = \exp(\tau_j) \langle \mathbf{s}_u^{(j)}, \mathbf{e}_i \rangle, \quad (16)$$

$\tau_j$  log-temperature  $\exp(\tau_j) \leq 20$  normalized state item vector  
 cosine similarity



□□ agent □□□□□ exponential moving-average baseline□

$$b_j \leftarrow 0.95b_j + 0.05\bar{r}_j, \quad A_j = r_j - b_j. \quad (26)$$

Policy-gradient □□□

$$\mathcal{L}_{\text{PG}} = - \sum_{j \in \{L, S, C\}} \mathbb{E} [A_j \log \pi_j(a_j | s_j)] - \beta \sum_j \mathbb{E} [H(\pi_j)], \quad (27)$$

□□ entropy coefficient  $\beta = 0.01$  □ Advantage □□□□□ detach □ gradient □□ log\_prob □  
□ logits □ LoRA □□□□□□□ sample() □

□□□ sampled-policy training □ joint stage □□□□ supervised anchor □

$$\mathcal{L}_{\text{sup}} = \frac{1}{3} \sum_{j \in \{L, S, C\}} \text{CE}(\mathbf{z}^{(j)}, 0). \quad (28)$$

□□□ long/short collapse □□□ specialization regularizer □

$$\mathcal{L}_{\text{spec}} = \cos^2(\mathbf{s}_u^{(L)}, \mathbf{s}_u^{(S)}). \quad (29)$$

□□ joint loss □

$$\mathcal{L}_{\text{joint}} = \mathcal{L}_{\text{PG}} + \lambda_{\text{sup}} \mathcal{L}_{\text{sup}} + \lambda_{\text{spec}} \mathcal{L}_{\text{spec}} \quad (30)$$

□□  $\lambda_{\text{sup}} = 0.1$  □  $\lambda_{\text{spec}} = 0.01$  □

## 6.4 Optimizer □ AMP □□□□□□

□□ stage □□□□ AdamW optimizer □□□ learning rates □□

| Stage       | Learning rate      | □□□□□                          |
|-------------|--------------------|--------------------------------|
| Specialists | $2 \times 10^{-4}$ | long □ short □ item projection |
| Coordinator | $2 \times 10^{-4}$ | coordinator                    |
| Joint RL    | $5 \times 10^{-5}$ | □ agent □ item projection      |

CUDA □□□□ FP16 autocast □ GradScaler □□□ optimizer step □□ global gradient norm □  
□□ 1.0 □□□□ joint learning rate □□□ RL □□□□ retrieval space □□□□

## 7 Gradient □ Credit Assignment

Coordinator loss □□□□□□□□ long/short states □ scores □ Joint stage □□□ gradient □□□  
□□□□ states □□□ specialist LoRA □

$$\frac{\partial \mathcal{L}_C}{\partial \theta_L} = \frac{\partial \mathcal{L}_C}{\partial \mathbf{z}^{\text{final}}} \frac{\partial \mathbf{z}^{\text{final}}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \theta_L}, \quad (31)$$

short agent □□□□□□□□ specialist □□□□□ sampled action □ reward □ policy loss □□□□  
□□□□ coordinator □□□□□□

□□ credit assignment □□ agent-local next-item reward □□□□□□□□□□□□□□□□  
□□□□□□□□□ COMA/QMIX □ counterfactual critic □□□□□□□ simulator □□□□□□  
□□

$$A_j^{\text{cf}} = Q(s, \mathbf{a}) - Q(s, \mathbf{a}_{-j}, a_j^{\text{baseline}}), \quad (32)$$

□□□□□□□ agent □□□ reward □□□□□□

## 8 □□□□□□□□□□

### 8.1 Full-catalog scoring

Evaluation □□□ 64-item sampled slate□□□□□□□□□□□□□□

$$E = W_p E_{\text{text}} \in \mathbb{R}^{N \times d}, \quad (33)$$

□□□□□□□□□□□□□□□□□□□□

$$\mathbf{z}_u = E \mathbf{s}_u. \quad (34)$$

□□□□□  $B \times N \times d$  □□□□□□□□□  $B \times d$  □  $d \times N$  □ GEMM□□□ catalog ranking □□□  
□□□□ GPU □□□□□

### 8.2 Seen-item masking

□ evaluation history □□□□□□ target □□□□□

$$z_{u,i} = -\infty, \quad (35)$$

□□□□□□□□□□□□ Validation history □ train□test history □ train □ valid□

### 8.3 Metrics

□□ target rank □□□□

$$\text{rank}_u = \sum_{i=1}^N \mathbb{I}[z_{u,i} \geq z_{u,i+}]. \quad (36)$$

□□□□□

$$\text{Recall@K} = \frac{1}{|U|} \sum_u \mathbb{I}[\text{rank}_u \leq K], \quad (37)$$

$$\text{NDCG@K} = \frac{1}{|U|} \sum_u \frac{\mathbb{I}[\text{rank}_u \leq K]}{\log_2(\text{rank}_u + 1)}, \quad (38)$$

□□  $K \in \{5, 10\}$  □□□□□ users/s □ user–item scores/s□□□□□□□□□□□□□□□□ catalog scoring □□□□

## 9 □□□□□□□□□□

1. □□□□□□□□□□ Mamba-2.8B □□□ dataset□item count □□□ fingerprint □□□ cache hit □□□□□□□□□□
2. □□ **base semantics**□□ agent □□□□□□□ projection□□□□□□□ 2.8B weights □□□ catalog embeddings□
3. □□ **history update**□specialist recurrence □□□□□□□□□□□□□□ memory □□□  $T^2$  attention map□



4. **history** 100 short agent 10 latency
5. 64-item slate policy evaluation GEMM catalog scoring
6. **Mixed precision** Mamba cache CUDA FP16 GradScaler gradient underflow
7. **Reservoir transition cap** 500,000 transitions Python object memory
8. **reason generator** checkpoint policy CPU CUDA cache

## 10

latency full-catalog scoring test users Top-1 item coordinator  $w_L, w_S$

Listing 1: prompt

```
Explain this recommendation in one concise sentence using only the
supplied history.
Long-term agent weight=<wL>; short-term agent weight=<wS>.
History: <recent item texts>.
Recommended item: <recommended item text>. Reason:
```

Mamba greedy decoding `do_sample=False` 40 token  
 agent weights recommendations.json post-hoc grounded explanation  
 reason quality RL reward reasoning-aware recommen-  
 dation item consistency history grounding reward

## 11

### 11.1 Launcher

```
# MovieLens
bash run_mamba_rl.sh movielens-1m 0
bash run_mamba_rl.sh movielens-20m 0

# Amazon Reviews 2023
bash run_mamba_rl.sh amazon-all-beauty 0
bash run_mamba_rl.sh amazon-movies-and-tv 0

# Yelp licensed local snapshots
bash run_mamba_rl.sh yelp19 0 /data/yelp-2019
bash run_mamba_rl.sh yelp23 0 /data/yelp-2023
```

Launcher DATASET CUDA\_ID DATA\_PATH MovieLens  
 Amazon Hugging Face cache Yelp

## 11.2 □□□□□□

Table 3: `run_mamba_rl.sh` □□□□□

| <span style="font-family: monospace;">□□</span> | <span style="font-family: monospace;">□□□</span> | <span style="font-family: monospace;">□□</span>                                                                                                                        |
|-------------------------------------------------|--------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| CACHE_DIR                                       | /workspace/P78123011/cache                       | HF <span style="font-family: monospace;">□□□□</span> Mamba vector cache <span style="font-family: monospace;">□</span>                                                 |
| BATCH_SIZE                                      | 128                                              | <span style="font-family: monospace;">□□□□□</span> batch size <span style="font-family: monospace;">□</span>                                                           |
| EVAL_BATCH_SIZE                                 | 128                                              | Full-catalog evaluation batch <span style="font-family: monospace;">□</span>                                                                                           |
| SPECIALIST_EPOCHS                               | 1                                                | Stage 1 epoch <span style="font-family: monospace;">□□</span>                                                                                                          |
| COORDINATOR_EPOCHS                              | 1                                                | Stage 2 epoch <span style="font-family: monospace;">□□</span>                                                                                                          |
| RL_EPOCHS                                       | 3                                                | Joint REINFORCE epoch <span style="font-family: monospace;">□□</span>                                                                                                  |
| MAX_TRANSITIONS                                 | 500000                                           | Prefix transitions <span style="font-family: monospace;">□□□</span>                                                                                                    |
| SEED                                            | 25252                                            | Python <span style="font-family: monospace;">□</span> NumPy <span style="font-family: monospace;">□</span> PyTorch seed <span style="font-family: monospace;">□</span> |

□□□

```
BATCH_SIZE=256 RL_EPOCHS=5 MAX_TRANSITIONS=1000000 \
  bash run_mamba_rl.sh movielens-20m 0
```

## 11.3 □□□ CPU smoke test

```
python -m src.train_mamba_rl \
  --dataset synthetic \
  --device cpu \
  --skip-mamba \
  --no-generate-reasons \
  --max-transitions 32 \
  --batch-size 8
```

```
--skip-mamba □□□□ seed □ random vectors □□□□□□□□□□□□□□□□□□□□□□□□
□□□
```

## 12 □□□□□□□□

□□ run □□□

```
outputs_mamba_rl/<dataset>/<YYYYMMDD_HHMMSS>/
|-- train.log
|-- multi_agent_lora.pt
|-- loss.png
|-- metrics.json
\-- recommendations.json
```

`multi_agent_lora.pt` □□□□ state dict□□□ CLI config□ validation/test metrics  
□ item-vector artifact path□□□ item feature buffer □□□□□ checkpoint□□□□□□□ artifact  
path □□□□□ cached item matrix□□□□ state dict□

**13**    □ □ □ □ □ □ □ □ □ □ □ □

- [illegible]

**14**    □ □ □ □ □ □ □

11

Table 4: Ablation study

| Configuration                             | Performance                    |
|-------------------------------------------|--------------------------------|
| agent vs. Multi-agent                     |                                |
| long+short+coordinator                    |                                |
| 10 window vs. 100/10 windows              |                                |
| 0.5/0.5 vs. learned coordinator           |                                |
| Stage 1/2 joint RL                        | variance collapse              |
| specialization loss                       | LoRA agents states             |
| supervised anchor                         | Pure REINFORCE                 |
| Mamba text vector vs. random/ID embedding | cold/sparse item               |
| Selective recurrence vs. GRU/-Transformer | history scaling latency memory |
| Sampled vs. hard negatives                | Full-catalog ranking gap       |
| reason generation                         | throughput                     |

Recall/NDCG catalog coverage novelty intra-list diversity history agent weight state cosine similarity training transitions/s evaluation users/s scores/s peak GPU memory reason latency

## 15 Ablation study

### 15.1 Reward-adaptive Mamba dynamics

$\delta_t$  feedback dwell time uncertainty

$$\delta_t = f_{\theta}(\mathbf{x}_t, r_{t-1}, \sigma_{t-1}), \quad (39)$$

reward-conditioned state dynamics

### 15.2 Offline RL

logged recommendation behavior policy propensity Conservative Q-Learning

$$\mathcal{L}_{\text{CQL}} = \log \sum_a \exp Q(s, a) - Q(s, a_{\text{logged}}), \quad (40)$$

out-of-distribution

### 15.3 Centralized critic simulator

KuaiSim RL4RS centralized training/decentralized execution critic agent actions session watch time retention diversity reward hit reward

## □□ policy □□□□ item □ reasoning trajectory□□□□

□□ group-relative policy optimization □□□□□□ trajectory □ gradient variance□

Multi-Agent Mamba-RL pipeline    Mamba    Mamba    selective state  
agent-specific LoRA    policy gradient    selective-state agents    coordinator    RL    GEMM    multi-agent recommendation    reward-adaptive dynamics    offline RL    centralized critic    reasoning reward